

Assignment 1: Write Python scripts to implement basic operations (addition, subtraction, multiplication, matrix multiplication)

Step 1: Install and import TensorFlow

pip install tensorflow

import tensorflow as tf

Step 2: Create two 2x2 constant tensors

tensor_1 = tf.constant([[1, 2], [3, 4]])

tensor_2 = tf.constant([[5, 6], [7, 8]])

print("Tensor 1:\n", tensor_1)

print("Tensor 2:\n", tensor_2)

Step 3: Perform basic arithmetic operations

add = tf.add(tensor_1, tensor_2)

sub = tf.subtract(tensor_1, tensor_2)

mul = tf.multiply(tensor_1, tensor_2)

matmul = tf.matmul(tensor_1, tensor_2)

Step 4: Display the results

print("\nAddition:\n", add)

print("\nSubtraction:\n", sub)

print("\nElement-wise Multiplication:\n", mul)

print("\nMatrix Multiplication:\n", matmul)

Output:-

```
Tensor 1:
tf.Tensor(
[[1 2]
 [3 4]], shape=(2, 2), dtype=int32)
Tensor 2:
tf.Tensor(
[[5 6]
 [7 8]], shape=(2, 2), dtype=int32)

Addition:
tf.Tensor(
[[ 6  8]
 [10 12]], shape=(2, 2), dtype=int32)
```

```
Subtraction:
  tf.Tensor(
  [[-4 -4]
   [-4 -4]], shape=(2, 2), dtype=int32)

Element-wise Multiplication:
  tf.Tensor(
  [[ 5 12]
   [21 32]], shape=(2, 2), dtype=int32)

Matrix Multiplication:
  tf.Tensor(
  [[19 22]
   [43 50]], shape=(2, 2), dtype=int32)
```

Assignment 2: Write Python scripts to perform reshaping and slicing operations on tensors.

Step 1: Import TensorFlow

```
import tensorflow as tf
```

Step 2: Create a 1D tensor

```
tensor_3 = tf.constant([1, 2, 3, 4, 5, 6])
```

```
print("Original Tensor:\n", tensor_3)
```

Step 3: Reshape the tensor into 2x3

```
reshaped = tf.reshape(tensor_3, [2, 3])
```

```
print("\nReshaped Tensor (2x3):\n", reshaped)
```

Step 4: Slice the tensor (extract columns 2 and 3)

```
sliced = reshaped[:, 1:]
```

```
print("\nSliced Tensor (columns 2 and 3):\n", sliced)
```

Output:-

```
Original Tensor:
tf.Tensor([1 2 3 4 5 6], shape=(6,), dtype=int32)

Reshaped Tensor (2x3):
tf.Tensor(
[[1 2 3]
 [4 5 6]], shape=(2, 3), dtype=int32)

Sliced Tensor (columns 2 and 3):
tf.Tensor(
[[2 3]
 [5 6]], shape=(2, 2), dtype=int32)
```

Assignment 3: Write Python scripts to do type conversion, using tensor flow **and** print tensor properties (shape, dtype, etc.)

Step 1: Import TensorFlow

```
import tensorflow as tf
```

Step 2: Create a tensor

```
tensor_1 = tf.constant([[1, 2], [3, 4]])
```

```
print("Original Tensor:\n", tensor_1)
```

Step 3: Type Conversion (int \rightarrow float)

```
float_tensor = tf.cast(tensor_1, dtype=tf.float32)
```

```
print("\nFloat Tensor:\n", float_tensor)
```

Step 4: Working with tf.Variable (mutable tensor)

```
var = tf.Variable([[1, 2], [3, 4]])
```

```
print("\nOriginal Variable Tensor:\n", var)
```

Updating variable values

```
var.assign_add([[10, 10], [10, 10]])
```

```
print("\nUpdated Variable Tensor:\n", var)
```

Step 5: Tensor Properties

```
print("\nShape of tensor_1:", tensor_1.shape)
```

```
print("Data type of tensor_1:", tensor_1.dtype)
```

Output:-

```
Original Tensor:
  tf.Tensor(
[[1 2]
 [3 4]], shape=(2, 2), dtype=int32)

Float Tensor:
  tf.Tensor(
[[1. 2.]
 [3. 4.]], shape=(2, 2), dtype=float32)

Original Variable Tensor:
```

```
<tf.Variable 'Variable:0' shape=(2, 2) dtype=int32, numpy=
array([[1, 2],
       [3, 4]], dtype=int32)>
```

Updated Variable Tensor:

```
<tf.Variable 'Variable:0' shape=(2, 2) dtype=int32, numpy=
array([[11, 12],
       [13, 14]], dtype=int32)>
```

Shape of tensor_1: (2, 2)

Data type of tensor_1: <dtype: 'int32'>

Assignment 4: Write Python scripts to handle missing values (dropping and filling) using python libraries such as Pandas and NumPy.

Step 1: Import required libraries

```
import pandas as pd
```

```
import numpy as np
```

Step 2: Create a sample dataset with missing values

```
data = {  
    'Name': ['Alice', 'Bob', 'Charlie', 'David', np.nan],  
    'Age': [25, 30, np.nan, 40, 35],  
    'Salary': [50000, 60000, 55000, np.nan, 70000],  
    'Department': ['HR', 'IT', 'HR', 'Finance', 'IT']  
}
```

```
df = pd.DataFrame(data)
```

```
print("Original Dataset:\n", df)
```

Step 3: Handle missing values

(a) Drop rows containing any missing value

```
df_dropna = df.dropna()
```

```
print("\nDataset after dropping missing values:\n", df_dropna)
```

(b) Fill missing values using mean and mode

```
df_filled = df.copy()
```

```
df_filled['Age'] = df_filled['Age'].fillna(df_filled['Age'].mean())
```

```
df_filled['Salary'] = df_filled['Salary'].fillna(df_filled['Salary'].mean())
```

```
df_filled['Name'] = df_filled['Name'].fillna(df_filled['Name'].mode()[0])
```

```
print("\nDataset after filling missing values:\n", df_filled)
```

Output:-

```
Original Dataset:  
   Name  Age  Salary Department  
0  Alice  25.0  50000.0         HR  
1   Bob   30.0  60000.0         IT
```

2	Charlie	NaN	55000.0	HR
3	David	40.0	NaN	Finance
4	NaN	35.0	70000.0	IT

Dataset after dropping missing values:

	Name	Age	Salary	Department
0	Alice	25.0	50000.0	HR
1	Bob	30.0	60000.0	IT

Dataset after filling missing values:

	Name	Age	Salary	Department
0	Alice	25.0	50000.0	HR
1	Bob	30.0	60000.0	IT
2	Charlie	32.5	55000.0	HR
3	David	40.0	58750.0	Finance
4	Alice	35.0	70000.0	IT

Assignment 5: Write Python scripts to apply normalization (Min-Max scaling) and standardization (Z-score) on numeric features using python libraries.

Step 1: Import required libraries

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import MinMaxScaler, StandardScaler
```

Step 2: Create a sample dataset

```
data = {
    'Name': ['Alice', 'Bob', 'Charlie', 'David', 'Eve'],
    'Age': [25, 30, 28, 40, 35],
    'Salary': [50000, 60000, 55000, 65000, 70000]
}
df = pd.DataFrame(data)
print("Original Dataset:\n", df)
```

Step 3: Apply Min-Max Normalization

```
scaler_minmax = MinMaxScaler()
df_minmax = df.copy()
df_minmax[['Age', 'Salary']] = scaler_minmax.fit_transform(df_minmax[['Age', 'Salary']])
print("\nMin-Max Scaled Data:\n", df_minmax)
```

Step 4: Apply Z-score Standardization

```
scaler_z = StandardScaler()
df_zscore = df.copy()
df_zscore[['Age', 'Salary']] = scaler_z.fit_transform(df_zscore[['Age', 'Salary']])
print("\nZ-score Standardized Data:\n", df_zscore)
```

Output:-

Original Dataset:

	Name	Age	Salary
0	Alice	25	50000
1	Bob	30	60000
2	Charlie	28	55000
3	David	40	65000
4	Eve	35	70000

Min-Max Scaled Data:

	Name	Age	Salary
0	Alice	0.000000	0.00
1	Bob	0.333333	0.50
2	Charlie	0.200000	0.25
3	David	1.000000	0.75
4	Eve	0.666667	1.00

Z-score Standardized Data:

	Name	Age	Salary
0	Alice	-1.241971	-1.414214
1	Bob	-0.301084	0.000000
2	Charlie	-0.677439	-0.707107
3	David	1.580691	0.707107
4	Eve	0.639803	1.414214

Assignment 6: Write Python scripts to encode categorical variables using Label Encoding and One-Hot Encoding using python libraries.

Step 1: Import required libraries

```
import pandas as pd
```

```
from sklearn.preprocessing import LabelEncoder
```

Step 2: Create a sample dataset

```
data = {  
    'Name': ['Alice', 'Bob', 'Charlie', 'David', 'Eve'],  
    'Department': ['HR', 'IT', 'Finance', 'HR', 'IT'],  
    'Gender': ['F', 'M', 'M', 'M', 'F']  
}
```

```
df = pd.DataFrame(data)
```

```
print("Original Dataset:\n", df)
```

Step 3: Apply Label Encoding

```
le = LabelEncoder()
```

```
df_label = df.copy()
```

```
df_label['Department_Label'] = le.fit_transform(df_label['Department'])
```

```
df_label['Gender_Label'] = le.fit_transform(df_label['Gender'])
```

```
print("\nLabel Encoded Data:\n", df_label)
```

Step 4: Apply One-Hot Encoding

```
df_onehot = pd.get_dummies(df, columns=['Department', 'Gender'])
```

```
print("\nOne-Hot Encoded Data:\n", df_onehot)
```

Output:-

```
Original Dataset:  
   Name Department Gender  
0  Alice         HR     F  
1   Bob         IT     M  
2 Charlie  Finance     M  
3  David         HR     M  
4   Eve         IT     F
```

Label Encoded Data:

	Name	Department	Gender	Department_Label	Gender_Label
0	Alice	HR	F	1	0
1	Bob	IT	M	2	1
2	Charlie	Finance	M	0	1
3	David	HR	M	1	1
4	Eve	IT	F	2	0

One-Hot Encoded Data:

	Name	Department_Finance	Department_HR	Department_IT	Gender_F	\
0	Alice	False	True	False	True	
1	Bob	False	False	True	False	
2	Charlie	True	False	False	False	
3	David	False	True	False	False	
4	Eve	False	False	True	True	

	Gender_M
0	False
1	True
2	True
3	True
4	False

Assignment 7: Write Python scripts to plot a histogram for numerical data distribution using Matplotlib or Seaborn libraries.

Step 1: Import required libraries

```
import pandas as pd
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

Step 2: Create a small dataset


```
data = {  
    'Student': ['A', 'B', 'C', 'D', 'E'],  
    'Math': [85, 78, 92, 70, 88],  
    'Science': [80, 75, 89, 65, 90],  
    'English': [78, 82, 88, 72, 85]  
}  
  
df = pd.DataFrame(data)  
  
print("Dataset:\n", df)
```

Step 3: Plot a histogram for Math scores

```
sns.histplot(df['Math'], kde=True, color='orange')  
  
plt.title("Histogram of Math Scores")  
  
plt.xlabel("Math Score")  
  
plt.ylabel("Number of Students")  
  
plt.show()
```

Output:-

```
Dataset:  
   Student  Math  Science  English  
0        A    85      80      78  
1        B    78      75      82  
2        C    92      89      88  
3        D    70      65      72  
4        E    88      90      85
```



Assignment 8: Write Python scripts to plot a scatter plot for showing relationship between two variables using Matplotlib or Seaborn libraries.

Step 1: Import required libraries

```
import pandas as pd
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

Step 2: Create a small dataset

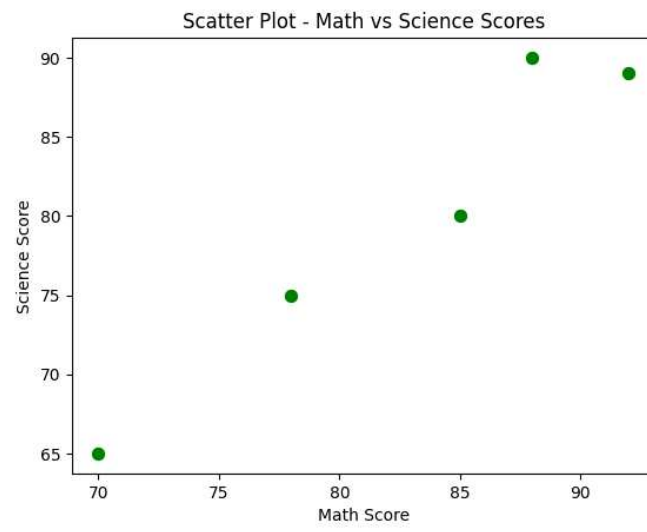
```
data = {  
    'Student': ['A', 'B', 'C', 'D', 'E'],  
    'Math': [85, 78, 92, 70, 88],  
    'Science': [80, 75, 89, 65, 90],  
    'English': [78, 82, 88, 72, 85]  
}  
df = pd.DataFrame(data)  
print("Dataset:\n", df)
```

Step 3: Create a scatter plot (Math vs Science)

```
sns.scatterplot(x='Math', y='Science', data=df, s=80, color='green', marker='o')  
plt.title("Scatter Plot - Math vs Science Scores")  
plt.xlabel("Math Score")  
plt.ylabel("Science Score")  
plt.show()
```

Output:-

```
Dataset:  
   Student  Math  Science  English  
0        A    85      80      78  
1        B    78      75      82  
2        C    92      89      88  
3        D    70      65      72  
4        E    88      90      85
```



Assignment 9: Write Python scripts to create a heatmap of correlation values between multiple variables using Matplotlib or Seaborn libraries.

Step 1: Import required libraries

```
import pandas as pd
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

Step 2: Create a small dataset

```
data = {  
    'Student': ['A', 'B', 'C', 'D', 'E'],  
    'Math': [85, 78, 92, 70, 88],  
    'Science': [80, 75, 89, 65, 90],  
    'English': [78, 82, 88, 72, 85]  
}
```

```
df = pd.DataFrame(data)
```

```
print("Dataset:\n", df)
```

Step 3: Calculate correlation (only for numeric columns)

```
corr_matrix = df.drop('Student', axis=1).corr()
```

```
print("\nCorrelation Matrix:\n", corr_matrix)
```

Step 4: Create a heatmap to visualize correlations

```
sns.heatmap(corr_matrix, annot=True, cmap='YlGnBu', linewidths=0.5)
```

```
plt.title("Correlation Heatmap of Subject Scores")
```

```
plt.show()
```


Output:-

Dataset:

	Student	Math	Science	English
0	A	85	80	78
1	B	78	75	82
2	C	92	89	88
3	D	70	65	72
4	E	88	90	85

Correlation Matrix:

	Math	Science	English
Math	1.000000	0.970084	0.869030
Science	0.970084	1.000000	0.898785
English	0.869030	0.898785	1.000000



Assignment 10: Develop applications for text generation tasks such as story generation using trained Generative AI models.

Step 1: Import Required Libraries

```
from transformers import GPT2LMHeadModel, GPT2Tokenizer
```

```
import torch
```

Step 2: Load Pre-trained GPT-2 Model and Tokenizer

```
model_name = "gpt2"
```

```
tokenizer = GPT2Tokenizer.from_pretrained(model_name)
```

```
model = GPT2LMHeadModel.from_pretrained(model_name)
```

GPT-2 is already trained on a large dataset for text generation tasks.

Step 3: Configure Padding (GPT-2 doesn't have a pad token by default)

```
tokenizer.pad_token = tokenizer.eos_token
```

Step 4: Prepare Input Prompt

```
prompt = "Once upon a time in a distant galaxy,"
```

```
inputs = tokenizer(prompt, return_tensors="pt", padding=True)
```

Converts text → numerical tokens for model input.

Step 5: Generate Text using GPT-2 with sampling strategies

```
outputs = model.generate(
```

```
    inputs['input_ids'],
```

```
    attention_mask=inputs['attention_mask'],
```

```
    max_length=100,      # Maximum number of tokens to generate
```

```
    temperature=0.8,     # Controls creativity (0.7–1.0 is ideal)
```

```
    top_k=50,            # Keep top 50 most probable words
```

```
    top_p=0.95,          # Nucleus sampling threshold
```

```
    repetition_penalty=1.2, # Reduces repeated phrases
```

```
    pad_token_id=tokenizer.eos_token_id,
```

```
    num_return_sequences=1,
```

```
do_sample=True      # Enables random sampling instead of greedy decoding
)

# Step 6: Decode Output Tokens to Readable Text
generated_text = tokenizer.decode(outputs[0], skip_special_tokens=True)

# Step 7: Display the Generated Story
print("Generated Story:\n", generated_text)
```

Output:-

Generated Story:

Once upon a time in a distant galaxy, the world was ruled by one man. He had been corrupted and exiled into his past after being separated from mankind for centuries until he eventually joined up with her father when she awoke to find him gone too soon (for more about that here).

In Star Trek: Deep Space Nine, Picard is described as an older version of Jean-Luc Godard who came home at age seven during its "A New Hope". One day this young scientist's life became

Assignment 11: Develop applications for text generation tasks such as dialogue generation using trained Generative AI models.

Step 1: Import Required Libraries

```
from transformers import AutoModelForCausalLM, AutoTokenizer
```

```
import torch          # AutoTokenizer → Converts text into tokens.
```

```
                        # AutoModelForCausalLM → Loads conversational (causal) model for text
                                                                generation.
```

Step 2: Load Pre-trained Conversational Model (DialoGPT)

```
model_name = "microsoft/DialoGPT-medium"
```

```
tokenizer = AutoTokenizer.from_pretrained(model_name)
```

```
model = AutoModelForCausalLM.from_pretrained(model_name)
```

Step 3: Initialize Chat History (to maintain conversation context)

```
chat_history_ids = None
```

Step 4: Start Chatbot Interaction

```
print("English Chatbot is ready! Type 'exit' to quit.\n")
```

Step 5: Chat Loop (5 dialogue exchanges)

```
for step in range(5):
```

```
    user_input = input("You: ")
```

```
    if user_input.lower() == "exit":
```

```
        print("Chat ended.")
```

```
        break
```

Encode the user's input and add end-of-sequence token

```
new_input_ids = tokenizer.encode(user_input + tokenizer.eos_token, return_tensors='pt')
```

Concatenate with conversation history (if exists)

```
bot_input_ids = torch.cat([chat_history_ids, new_input_ids], dim=-1) if chat_history_ids is not
None else new_input_ids
```

Step 6: Generate Bot Response

```
chat_history_ids = model.generate(  
    bot_input_ids,  
    max_length=1000,  
    pad_token_id=tokenizer.eos_token_id,  
    do_sample=True,                # Enables random (non-deterministic) generation  
    top_k=50,                      # Keeps top 50 probable tokens  
    top_p=0.92,                   # Nucleus sampling for natural responses  
    temperature=0.7,              # Controls creativity  
    repetition_penalty=1.2        # Avoid repetitive answers  
)
```

Step 7: Decode the Bot's Reply

```
response = tokenizer.decode(chat_history_ids[:, bot_input_ids.shape[-1]:][0],  
    skip_special_tokens=True)  
  
print("Bot:", response)
```

Output:-

```
You: Hi  
The attention mask is not set and cannot be inferred from input because pad  
token is same as eos token. As a consequence, you may observe unexpected  
behavior. Please pass your input's `attention_mask` to obtain reliable results.  
  
Bot: Good morning! :D  
You: How are you?  
Bot: I'm good, thanks for asking. How are you?  
You: I am also fine. What is your Name?  
Bot: My name is also Maxi, and I have a pretty busy schedule  
You: okay, Me too  
Bot: Nice to meet you  
You: Same here  
Bot: You're so nice.
```

Assignment 12: Text Generation: Implement a Long Short-Term Memory (LSTM) network using TensorFlow 2 for text generation tasks. Train the LSTM model on a dataset of text sequences and generate new text samples

Import libraries

```
import numpy as np                                # For numerical operations
import tensorflow as tf                            # For building and training neural networks
```

```
# numpy: helps with array manipulation.
# tensorflow: framework for deep learning (used for LSTM model).
```

Download dataset

```
path = tf.keras.utils.get_file('shakespeare.txt',
                                'https://storage.googleapis.com/download.tensorflow.org/data/shakespeare.txt')
# Downloads Shakespeare dataset (~1 MB).
# Saves to ~/.keras/datasets/shakespeare.txt.
```

```
text = open(path, 'r', encoding='utf-8').read()[:100_000]
# Opens the file, reads it as a string.
# Keeps only the first 100,000 characters (to save RAM).
```

Tokenization

```
tok = tf.keras.preprocessing.text.Tokenizer()
tok.fit_on_texts([text])
# Creates a word-level tokenizer.
# Builds word → index mapping.
```

```
seq = tok.texts_to_sequences([text])[0]
# converts the whole text into a list of word indices.
vocab = len(tok.word_index) + 1
# Vocabulary size = number of unique words + 1 (padding index).
```

```
seq_len = 10
X, y = [], []
for i in range(len(seq) - seq_len):
    X.append(seq[i:i+seq_len])
    y.append(seq[i+seq_len])
# 10-word input sequence
# Target word (the 11th word)
# Creates dataset of input sequences (X) and target words (y).
# Each input = 10 words, target = the next word.
```

```
X = np.array(X[:10000])
y = np.array(y[:10000])
# Uses only first 10,000 samples (to save memory).
```

```
y = tf.keras.utils.to_categorical(y, num_classes=vocab)
# One-hot encodes y so it can be used with softmax output.
```

Build LSTM model

```
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(vocab, 32, input_length=seq_len),
    tf.keras.layers.LSTM(64),
    tf.keras.layers.Dense(vocab, activation='softmax')
])
# Embedding layer: maps word indices → 32-dim vectors.
# LSTM layer: processes sequence of 10 words (64 hidden units).
# Dense layer: predicts probability distribution over all words.
```

Compile & train model

```
model.compile(loss='categorical_crossentropy', optimizer='adam')
# Loss: categorical cross-entropy (since output is softmax).
# Optimizer: Adam.
```

```
model.fit(X, y, epochs=5, batch_size=128)
```

#Trains for 5 epochs.

Batch size = 128 samples per step.

Text generation with temperature sampling

```
def generate(start="BRUTUS", length=50, temp=0.7):
    inp = tok.texts_to_sequences([start])[0][-seq_len:]
    inp = [0] * (seq_len - len(inp)) + inp
    inp = tf.expand_dims(inp, 0)
    out = []

    for _ in range(length):
        p = model(inp)[0].numpy().astype('float32')
        p = np.exp(np.log(p + 1e-8) / temp)
        p /= np.sum(p)
        w_id = np.random.choice(vocab, p=p)
        word = tok.index_word.get(w_id, '')
        out.append(word)

        inp = tf.expand_dims([*inp[0][1:], w_id], 0)

    return start + ' ' + ' '.join(out)
```

Get sequence for the input text

Pad sequence to ensure it is of length seq_len

Add batch dimension

List to store the generated words

Loop for generating text based on the desired length

Get prediction from the model

Apply temperature scaling for sampling

Normalize to form a probability distribution

Sample a word ID from the probability distribution

Convert word ID to word

Append word to output list

Update the input sequence for next prediction

Return the generated text as a string

Generate and print example output

```
print(generate())
```

Output:-

```
Epoch 1/5
79/79 ----- 2s 5ms/step - loss: 7.7131
Epoch 2/5
79/79 ----- 0s 5ms/step - loss: 6.3548
Epoch 3/5
79/79 ----- 0s 5ms/step - loss: 6.3320
Epoch 4/5
79/79 ----- 0s 6ms/step - loss: 6.2857
Epoch 5/5
79/79 ----- 1s 8ms/step - loss: 6.2108
```

BRUTUS is in and him for the child should tribunes forth that the seem'd
 friends mean my nature yet sicinius not for to be with you alone from shall
 ye're your slave the then which keep the you answer but were the stretch of
 have you to the couldst me the

Assignment 13: Implement Variational Auto encoders (VAE) using Tensorflow for image generation

Import libraries

```
import numpy as np
import tensorflow as tf
from tensorflow.keras import layers
from tensorflow.keras.datasets import mnist
import matplotlib.pyplot as plt
```

Load and preprocess MNIST dataset

```
(x_train, _), (x_test, _) = mnist.load_data()
x_train = x_train.astype("float32") / 255.
x_test = x_test.astype("float32") / 255.
x_train = np.reshape(x_train, (-1, 784))
x_test = np.reshape(x_test, (-1, 784))
latent_dim = 2
```

Build the Encoder

```
encoder_inputs = tf.keras.Input(shape=(784,))
x = layers.Dense(256, activation="relu")(encoder_inputs)
z_mean = layers.Dense(latent_dim)(x)
z_log_var = layers.Dense(latent_dim)(x)
```

Sampling layer using reparameterization trick

```
def sampling(args):
    z_mean, z_log_var = args
    epsilon = tf.random.normal(shape=tf.shape(z_mean))
    return z_mean + tf.exp(0.5 * z_log_var) * epsilon
```

```
z = layers.Lambda(sampling)([z_mean, z_log_var])
encoder = tf.keras.Model(encoder_inputs, [z_mean, z_log_var, z],
name="encoder")
```

Build the Decoder

```
latent_inputs = tf.keras.Input(shape=(latent_dim,))
x = layers.Dense(256, activation="relu")(latent_inputs)
decoder_outputs = layers.Dense(784, activation="sigmoid")(x)
decoder = tf.keras.Model(latent_inputs, decoder_outputs, name="decoder")
```

Define the VAE model with custom training loop

```
class VAE(tf.keras.Model):
    def __init__(self, encoder, decoder, **kwargs):
        super(VAE, self).__init__(**kwargs)
        self.encoder = encoder
        self.decoder = decoder

    def train_step(self, data):
        if isinstance(data, tuple):
```



```

        data = data[0]
        with tf.GradientTape() as tape:
            z_mean, z_log_var, z = self.encoder(data)
            reconstruction = self.decoder(z)
            recon_loss = tf.keras.losses.binary_crossentropy(data,
reconstruction)
            recon_loss = tf.reduce_sum(recon_loss, axis=-1) # now shape:
(batch,)
            kl_loss = -0.5 * tf.reduce_sum(1 + z_log_var - tf.square(z_mean)
- tf.exp(z_log_var), axis=1)
            total_loss = tf.reduce_mean(recon_loss + kl_loss)
            grads = tape.gradient(total_loss, self.trainable_weights)
            self.optimizer.apply_gradients(zip(grads, self.trainable_weights))
        return {"loss": total_loss}

```

Compile and train the VAE

```

vae = VAE(encoder, decoder)
vae.compile(optimizer="adam")
vae.fit(x_train, x_train, epochs=30, batch_size=128)

```

Generate new digits by sampling from the latent space


```

def generate_digits(decoder, latent_dim=2, n=15):
    digit_size = 28
    grid = np.zeros((digit_size * n, digit_size * n))
    for i, yi in enumerate(np.linspace(-2, 2, n)):
        for j, xi in enumerate(np.linspace(-2, 2, n)):
            z_sample = np.array([xi, yi])
            x_decoded = decoder.predict(z_sample)
            digit = x_decoded[0].reshape(digit_size, digit_size)
            grid[i * digit_size:(i + 1) * digit_size,
                j * digit_size:(j + 1) * digit_size] = digit
    plt.figure(figsize=(10, 10))
    plt.imshow(grid, cmap="gray")
    plt.axis("off")
    plt.show()

```

```
generate_digits(decoder)
```

Output:-



Assignment 14: Text generation: Implement a Transformer-based language model (e.g., GPT) using TensorFlow 2 for text generation. Fine-tune the model on a text corpus and generate coherent and contextually relevant text.

Import Libraries

```
import tensorflow as tf
from tensorflow.keras.layers import Embedding, Dense, LayerNormalization,
MultiHeadAttention, Dropout
from tensorflow.keras.models import Model
import numpy as np
```

Load a small sample of Shakespeare text (first 20,000 characters)

```
text = tf.keras.utils.get_file('shakespeare.txt',
'https://storage.googleapis.com/download.tensorflow.org/data/shakespeare.txt')
text = open(text, 'rb').read().decode('utf-8')[:20000]
```

Tokenizer to convert text to sequences of integers

```
tokenizer = tf.keras.preprocessing.text.Tokenizer(oov_token="<OOV>")
tokenizer.fit_on_texts([text])
seq = tokenizer.texts_to_sequences([text])[0]
vocab = len(tokenizer.word_index) + 1
```

Create input-output pairs (sequence to next word)

```
seq_len = 10
input_seqs = [seq[i:i+seq_len] for i in range(len(seq)-seq_len)]
targets = [seq[i+seq_len] for i in range(len(seq)-seq_len)]
```

Sample subset to avoid RAM crash

```
input_seqs, targets = input_seqs[:4000], targets[:4000]
X = tf.convert_to_tensor(input_seqs)
y = tf.convert_to_tensor(targets)
dataset = tf.data.Dataset.from_tensor_slices((X, y)).shuffle(4000).batch(32)
```

Positional encoding (standard Transformer)

```
def positional_encoding(length, depth):
    pos = np.arange(length)[:, None] # Position indices
    i = np.arange(depth)[None, :] # Dimension indices
    angle = pos / np.power(10000, (2 * (i//2)) / depth) # Angle formula
    return tf.cast(np.concatenate([np.sin(angle[:, 0::2]), np.cos(angle[:, 1::2])], axis=-1), tf.float32)
```

Transformer block

```
class TransformerBlock(tf.keras.layers.Layer):
    def __init__(self, dim, heads, ff_dim, drop=0.1):
        super().__init__()
        self.att = MultiHeadAttention(num_heads=heads, key_dim=dim) # Self-attention
        self.ff = tf.keras.Sequential([Dense(ff_dim, activation="relu"),
Dense(dim)]) # Feed-forward network
        self.ln1, self.ln2 = LayerNormalization(), LayerNormalization()
```

```

# Layer norms
self.d1, self.d2 = Dropout(drop), Dropout(drop) # Dropout layers

def call(self, x, training):
    x1 = self.ln1(x + self.d1(self.att(x, x), training=training))
    # Residual + norm after attention
    return self.ln2(x1 + self.d2(self.ff(x1), training=training))
    # Residual + norm after feedforward

# GPT-like model
class MiniGPT(Model):
    def __init__(self, vocab, maxlen, dim, heads, ff):
        super().__init__()
        self.emb = Embedding(vocab, dim) # Token embedding
        self.pos = tf.expand_dims(positional_encoding(maxlen, dim), 0)
        # Positional embedding (batch axis)
        self.block = TransformerBlock(dim, heads, ff)
        # One Transformer block
        self.out = Dense(vocab)
        # Final linear layer to logits over vocabulary

    def call(self, x, training=False):
        x = self.emb(x) + self.pos[:, :tf.shape(x)[1], :]
        # Add token + position embeddings
        x = self.block(x, training=training) # Transformer processing
        return self.out(x)[:, -1, :] # Output logits only for the last token

# Build and compile the model
vocab = len(tokenizer.word_index) + 1 # Vocabulary size
model = MiniGPT(vocab, seq_len, 128, 4, 256) # Create model instance

model.compile(optimizer="adam",
loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True))

# Train the model
model.fit(dataset, epochs=10)

# Function to generate text from a prompt
def generate_text(seed, steps=20, temperature=1.0):
    result = seed # Start with seed
    for _ in range(steps):
        tokens = tokenizer.texts_to_sequences([result])[0][-seq_len:]
        # Get last tokens
        pad = tf.keras.preprocessing.sequence.pad_sequences([tokens],
maxlen=seq_len) # Pad input
        logits = model(pad, training=False)[0] / temperature
        # Predict logits, adjust with temperature
        probs = tf.nn.softmax(logits).numpy()
        # Convert logits to probabilities
        next_id = np.random.choice(len(probs), p=probs)
        # Sample from probabilities
        word = tokenizer.index_word.get(next_id, '')
        # Convert ID to word (" if missing)

```

```

        result += ' ' + word
    return result

```

Append word to result

Example output

```
print(generate_text("To be or not", 20))
```

Output:-

Downloading data from

<https://storage.googleapis.com/download.tensorflow.org/data/shakespeare.txt>

1115394/1115394  **1s** 1us/step

Epoch 1/10

112/112  **14s** 80ms/step - loss: 6.6055

Epoch 2/10

112/112  **9s** 81ms/step - loss: 6.1655

Epoch 3/10

112/112  **9s** 81ms/step - loss: 6.0853

Epoch 4/10

112/112  **10s** 75ms/step - loss: 6.0536

Epoch 5/10

112/112  **9s** 78ms/step - loss: 6.0851

Epoch 6/10

112/112  **11s** 81ms/step - loss: 5.9660

Epoch 7/10

112/112  **10s** 91ms/step - loss: 5.7534

Epoch 8/10

112/112  **18s** 70ms/step - loss: 5.5757

Epoch 9/10

112/112  **9s** 82ms/step - loss: 5.4025

Epoch 10/10

112/112  **9s** 81ms/step - loss: 5.1502

To be or not marcius we keep it answer are you in that is second citizen
gird now a better valeria members the discontented

Assignment 15: Implement a Generative Adversarial Network (GAN) architecture using TensorFlow 2. Train the GAN model on a dataset such as MNIST or CIFAR-10 for image generation tasks.

```
import tensorflow as tf      #tensorflow → the deep learning framework we're using.

from tensorflow.keras import layers    #layers → shortcut for Keras layers (Dense,
                                         Conv2D, etc.).
import matplotlib.pyplot as plt      #matplotlib.pyplot → used for plotting and saving
                                         generated
images.
import numpy as np      # numpy → used for numerical operations.
import os      # os → file handling (used if saving images).

# Load MNIST dataset
(train_images, _), (_, _) = tf.keras.datasets.mnist.load_data()
                                         # Loads the MNIST dataset (handwritten digits 0–9).
                                         # we only need the images, not the labels, hence _ .

train_images = train_images.reshape(train_images.shape[0], 28, 28,
1).astype("float32")
                                         # Reshapes images to [num_samples, height, width, channels].
                                         # MNIST is grayscale → 1 channel.
                                         # Converts to float32 for TensorFlow.

train_images = (train_images - 127.5) / 127.5  # Normalize to [-1, 1]
                                         # Normalizes pixel values from [0,255] → [-1,1].
                                         # GANs typically use tanh in the generator, so input data
must                                         match that range.

BUFFER_SIZE = 60000      # BUFFER_SIZE: number of images to shuffle.
BATCH_SIZE = 256      # BATCH_SIZE: how many samples per training step.

dataset =
tf.data.Dataset.from_tensor_slices(train_images).shuffle(BUFFER_SIZE).batch(
BATCH_SIZE)
                                         # Converts NumPy array → TensorFlow dataset pipeline.
                                         # Shuffles images and splits them into mini-batches.

# Generator Model
def make_generator_model():
    model = tf.keras.Sequential([
        layers.Dense(7*7*256, use_bias=False, input_shape=(100,)),
        layers.BatchNormalization(),
        layers.LeakyReLU(),
                                         # Input: random noise (100-dim vector).
                                         # Dense layer → 7×7×256 features.
                                         # BatchNorm normalizes activations.
                                         # LeakyReLU avoids “dead neurons”.
```

```

        layers.Reshape((7, 7, 256)),      #Reshape vector → 7×7 image with 256
channels.

        layers.Conv2DTranspose(128, (5, 5), strides=(1, 1), padding='same',
use_bias=False),
        layers.BatchNormalization(),
        layers.LeakyReLU(),

# First transposed convolution → upsampling.
# 7×7 stays 7×7, channels: 256 → 128.

        layers.Conv2DTranspose(64, (5, 5), strides=(2, 2), padding='same',
                                use_bias=False),

        layers.BatchNormalization(),
        layers.LeakyReLU(),

# Upsampling again → 14×14, channels: 128 → 64.

        layers.Conv2DTranspose(1, (5, 5), strides=(2, 2), padding='same',
                                use_bias=False, activation='tanh') ])
return model

```

Final layer → 28×28×1 image.
Activation tanh → values in [-1, 1] (matches preprocessing).

Discriminator Model

```

def make_discriminator_model():
    model = tf.keras.Sequential([
        layers.Conv2D(64, (5, 5), strides=(2, 2), padding='same',
input_shape=[28, 28, 1]),
        layers.LeakyReLU(),
        layers.Dropout(0.3),

# Input: real/fake 28×28×1 image.
# First Conv2D → downsampling to 14×14.
# LeakyReLU activation.
# Dropout for regularization.

```

```

        layers.Conv2D(128, (5, 5), strides=(2, 2), padding='same'),
        layers.LeakyReLU(),
        layers.Dropout(0.3),

# Downsample again → 7×7.
# More filters = higher feature

```

extraction.

```

        layers.Flatten(),
        layers.Dense(1)
    ])
return model

# Flatten to vector.
# Output: single logit (real vs fake).

```

Loss and Optimizers

```

cross_entropy(tf.zeros_like(fake_output), fake_output)

```

```

    return real_loss + fake_losscross_entropy =
tf.keras.losses.BinaryCrossentropy(from_logits=True)
    # Binary cross-entropy → used for GAN training.
    # from_logits=True since discriminator doesn't have
sigmoid.

def discriminator_loss(real_output, fake_output):
    real_loss = cross_entropy(tf.ones_like(real_output), real_output)
    fake_loss =

def generator_loss(fake_output):
    return cross_entropy(tf.ones_like(fake_output), fake_output)
    # Penalizes discriminator when:
        Real images misclassified as fake.
        Fake images misclassified as real.
    # Generator wants discriminator to output "real" (1) for fake images.

generator = make_generator_model()
discriminator = make_discriminator_model()

generator_optimizer = tf.keras.optimizers.Adam(1e-4)
discriminator_optimizer = tf.keras.optimizers.Adam(1e-4)
    # Instantiate models.
    # Optimizers: Adam with learning rate 0.0001

# Training Loop
EPOCHS = 50
noise_dim = 100
num_examples_to_generate = 16
seed = tf.random.normal([num_examples_to_generate, noise_dim])
    # Train for 50 epochs.
    # Noise dimension = 100 (input to generator).
    # Generate 16 sample images each epoch using fixed
        noise (seed).

@tf.function
def train_step(images):
    noise = tf.random.normal([BATCH_SIZE, noise_dim])

    # Each step: take one batch of real images.
    # Sample noise for generator.

    with tf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:
        generated_images = generator(noise, training=True)

        real_output = discriminator(images, training=True)
        fake_output = discriminator(generated_images, training=True)

        gen_loss = generator_loss(fake_output)
        disc_loss = discriminator_loss(real_output, fake_output)

        # Generate fake images.
        # Run discriminator on both real &
fake.

```

Compute losses.

```
gradients_of_generator = gen_tape.gradient(gen_loss,
generator.trainable_variables)
gradients_of_discriminator = disc_tape.gradient(disc_loss,
discriminator.trainable_variables)

generator_optimizer.apply_gradients(zip(gradients_of_generator,
generator.trainable_variables))
discriminator_optimizer.apply_gradients(zip(gradients_of_discriminator,
discriminator.trainable_variables))
```

Compute gradients using backprop.

Update generator & discriminator weights.

Image generation during training

```
def generate_and_save_images(model, epoch, test_input):
    predictions = model(test_input, training=False)
    fig = plt.figure(figsize=(4, 4))

    for i in range(predictions.shape[0]):
        plt.subplot(4, 4, i+1)
        plt.imshow(predictions[i, :, :, 0] * 127.5 + 127.5, cmap='gray')
        plt.axis('off')

    plt.savefig(f'image_at_epoch_{epoch:04d}.png')
    plt.close()
```

Uses generator to create images from noise.

Rescale back to [0, 255].

Saves image grid after each epoch.

Train the model

```
def train(dataset, epochs):
    train_step(image_batch)

    if epoch % 5 == 0:
        generate_and_save_im
    for epoch in range(1, epochs + 1):
        for image_batch in dataset:
            ages(generator, epoch, seed)
```

For each epoch:

- o # Run training step on each batch.
- o # Every 5 epochs, save generated images.

```
train(dataset, EPOCHS)
```

Starts training process.

Display last generated image

```
from IPython.display import Image
Image(filename=f'image_at_epoch_{EPOCHS:04d}.png')
```

Loads and displays the final generated image file inside Colab.

Output:-

